

# node.js on a PaaS

## the awesome and the wonky

Patrick Mueller - [@pmuellr](https://twitter.com/pmuellr) - [muellerware.org](http://muellerware.org)

developer advocate for IBM's [Bluemix](https://www.ibm.com/cloud/ibm-bluemix) PaaS

<http://pmuellr.github.io/slides/2014/10-node-js-on-a-paas>

<http://pmuellr.github.io/slides/> (all slides)

# wat - node.js on a PaaS

- you know what node.js is
- PaaS == Platform as a Service
- examples:
  - Heroku
  - Nodejitsu
  - Cloud Foundry
  - OpenShift
  - many more!

# PaaS fundamentals

- OS provided for you (Linux)
- you provide the application (web server)
- configure:
  - number of instances running
  - RAM per instance
  - ephemeral disk per instance
- run a few commands

voila!

web server running on the "cloud" (public internet)

# PaaS usage scenarios

- special focus on web servers, typically with:
  - one open HTTP port open when app starts
  - HTTPS support
  - free domain, or use your own custom domain
  - WebSocket support
- or anything - arbitrary compute

# PaaS app development methodology

The Twelve Factor App - <http://12factor.net/>

*patterns for building apps on the cloud*

- |                        |                     |
|------------------------|---------------------|
| 1. Codebase            | 7. Port binding     |
| 2. Dependencies        | 8. Concurrency      |
| 3. Config              | 9. Disposability    |
| 4. Backing Services    | 10. Dev/prod parity |
| 5. Build, release, run | 11. Logs            |
| 6. Processes           | 12. Admin processes |

# PaaS app development methodology

- **Codebase** - One codebase tracked in revision control, many deploys
- **Dependencies** - Explicitly declare and isolate dependencies
- **Config** - Store config in the environment
- **Backing Services** - Treat backing services as attached resources
- **Build, release, run** - Strictly separate build and run stages
- **Processes** - Execute the app as one or more stateless processes

# PaaS app development methodology

- **Port binding** - Export services via port binding
- **Concurrency** - Scale out via the process model
- **Disposability** - Maximize robustness with fast startup and graceful shutdown
- **Dev/prod parity** - Keep development, staging, and production as similar as possible
- **Logs** - Treat logs as event streams
- **Admin processes** - Run admin/management tasks as one-off processes

# let's deploy an app to Cloud Foundry

```
$ git clone https://github.com/pmuellr/cf-node-hello.git
$ cd cf-node-hello
$ cf push
Using manifest file ../manifest.yml
...
Installing IBM SDK for Node.js from admin cache
...
Installing dependencies
...
Uploading droplet (8.0M)
...
App started
...
urls: cf-node-hello-pjm.mybluemix.net
$ curl https://cf-node-hello-pjm.mybluemix.net
Hello World
```



## what just happened?

- **cf push** uploaded your application files to a staging server
- staging server got node binaries, package dependencies, packaged into archive
- vm allocated to run the app, archive downloaded, expanded, started
- web server now running on the internet

# how the staging server "builds" an app

- driven from **package.json**
- get node executable from **engines.node**
  - **{ "engines" : { "node" : "0.10.x" } }**
- **npm install** will be run to obtain packages
- for Heroku and Cloud Foundry, the build is scripted with a "buildpack"; you can also write your own

# PaaS tools

- web dashboard
- command-line tooling
  - Heroku toolbelt - **heroku**
  - Cloud Foundry - **cf**
- typically will be using both web and cli

# domains

- typically PaaS provides you a free domain for your apps:
  - **foo.herokuapp.com**
  - **bar.mybluemix.net**
  - **yow.cfapps.io**
- host names must be unique across free domain!
- custom domains usually supported

# https support

- most PaaS's provide SSL termination
- allows you to support http and https traffic with just an http server
- or do you want https all the way to your server?
- https support for custom domains not simple; upload certs, etc

## using hosted services

Instead of running your own database, queueing server, etc, you'll be using hosted services, like:

- MongoLab
- Cloudant
- Redis Cloud

Some PaaS's co-locate hosted services in same datacenter to reduce latency.

## using hosted services

- add service to your app via:
  - command-line tool
  - web dashboard
  - roll your own access any 3rd service however you can
- services exposed to app via environment variables
  - Heroku's **MONGODB\_URL** env var
  - inside of Cloud Foundry's **VCAP\_SERVICES** env var

## using hosted services

Heroku:

```
$ heroku addons:add mongolab
```

Cloud Foundry:

```
$ cf create-service mongolab sandbox my-mongo-db
```



# adapting your app

- configuration provided via environment variables
- **process.env** is your new best friend
- **process.env.PORT** - env var set to the port to bind server
- using Heroku's MongoLab add-on service:

```
process.env.MONGODB_URL // MongoDB db URL
// mongodb://[user]:[pass]@[host]:[port]/[path]
```

# adapting your app - Cloud Foundry

- **VCAP\_SERVICES** - JSON string containing structured service info
- **VCAP\_APPLICATION** - JSON string containing other environmental info
  - url(s) to server
  - ip address of server
  - start time
  - etc

# adapting your app - Cloud Foundry

VCAP\_SERVICES will look like this, but with even more data:

```
{
  "mongodb-2.2": [
    {
      "name": "mongodb",
      "label": "mongodb-2.2",
      "credentials": {
        "url": "mongodb://..."
      }
    }
  ]
}
```

# adapting your app - Cloud Foundry

use the [cfenv package](#) to programmatically introspect over **VCAP\_SERVICES** and **VCAP\_APPLICATION**

instead of:

```
services = JSON.parse(process.env.VCAP_SERVICES)
mongoURL = services["mongodb-2.2"][0].credentials.url
```

use this:

```
cfenv = require("cfenv")
appEnv = cfenv.getAppEnv()
mongoURL = appEnv.getServiceURL("mongodb")
```

# scaling

By default, PaaS will run one instance of your app. Want more?

Heroku:

```
$ heroku ps:scale web=42
```

Cloud Foundry:

```
$ cf scale my-app -i 42
```

# scaling

if you want to scale, servers must be stateless

- no caching mutable data
- sometimes you want to scale **down**, so be prepared for servers to end
  - no long running, non-atomic transactions

## **the awesome: summary**

- no special "cloud libraries" for your app to use
- quick deploy of applications to the cloud
- don't need to worry about installing operating systems
- don't need to worry about installing services
- easy to scale up/down

# **the wonky**

## **and how to de-wonk-ify**



## core issues

- can't configure base operating system
- often no ssh
- stdout/err via syslog
- ephemeral file system
  - database all the things

# debugging

node-inspector difficult to run

- wants two ports open
- PaaS typically provides only one

=>

use a proxy splitter

- run two servers, split traffic by URL
- [cf-node-debug](#)

## other diagnostic tools

- New Relic
  - get account, apply light config
  - instrument via **require("newrelic")**
- StrongLoop
  - requires StrongLoop's tooling
- PaaS-specific
  - Bluemix Monitoring and Analytics

# diagnosing startup problems

- app runs fine on your laptop
- fails when run on PaaS

=>

- most errors here are in startup
- at startup:
  - add LOTS of logging
  - add LOTS of error checking

## private packages

if the PaaS runs **npm install** for you, how do can you access private packages

=>

- manage them separately from the rest of your packages
  - see [pmuellr/bluemix-private-packages](#) for an example project structure
- arrange to use a private package manager (npm in the future)

# logging

- typically only get last **XX** lines of your stdout/err
- but easy hook-ups to logging services:
  - Loggly
  - PaperTrail
  - `splunk>storm`
  - SumoLogic
  - Logentries
  - roll your own syslog support

# dependency versions

```
{  
  "name": "my-awesome-app",  
  "dependencies": {  
    "express": "*"   
  }  
}
```

What's wrong with this **package.json** file?

# dependency versions

```
{  
  "name": "my-awesome-app",  
  "dependencies": {  
    "express": "*"   
  }  
}
```

Guess what happened the day [express 4.0](#) was released?

- removed from express:
  - properties on **express**, **req**, **res**
  - all bundled middleware except static



# dependency versions

Lesson: lock down your dependencies

do one of these:

- version/put your node\_modules
- use fixed major/minor version specs: eg, "**3.4.x**"
- use [npm shrinkwrap](#)

**fin**